


BUILDING

**SCALABLE +
MAINTAINABLE +
RELIABLE**

APPS


12-factors to consider...





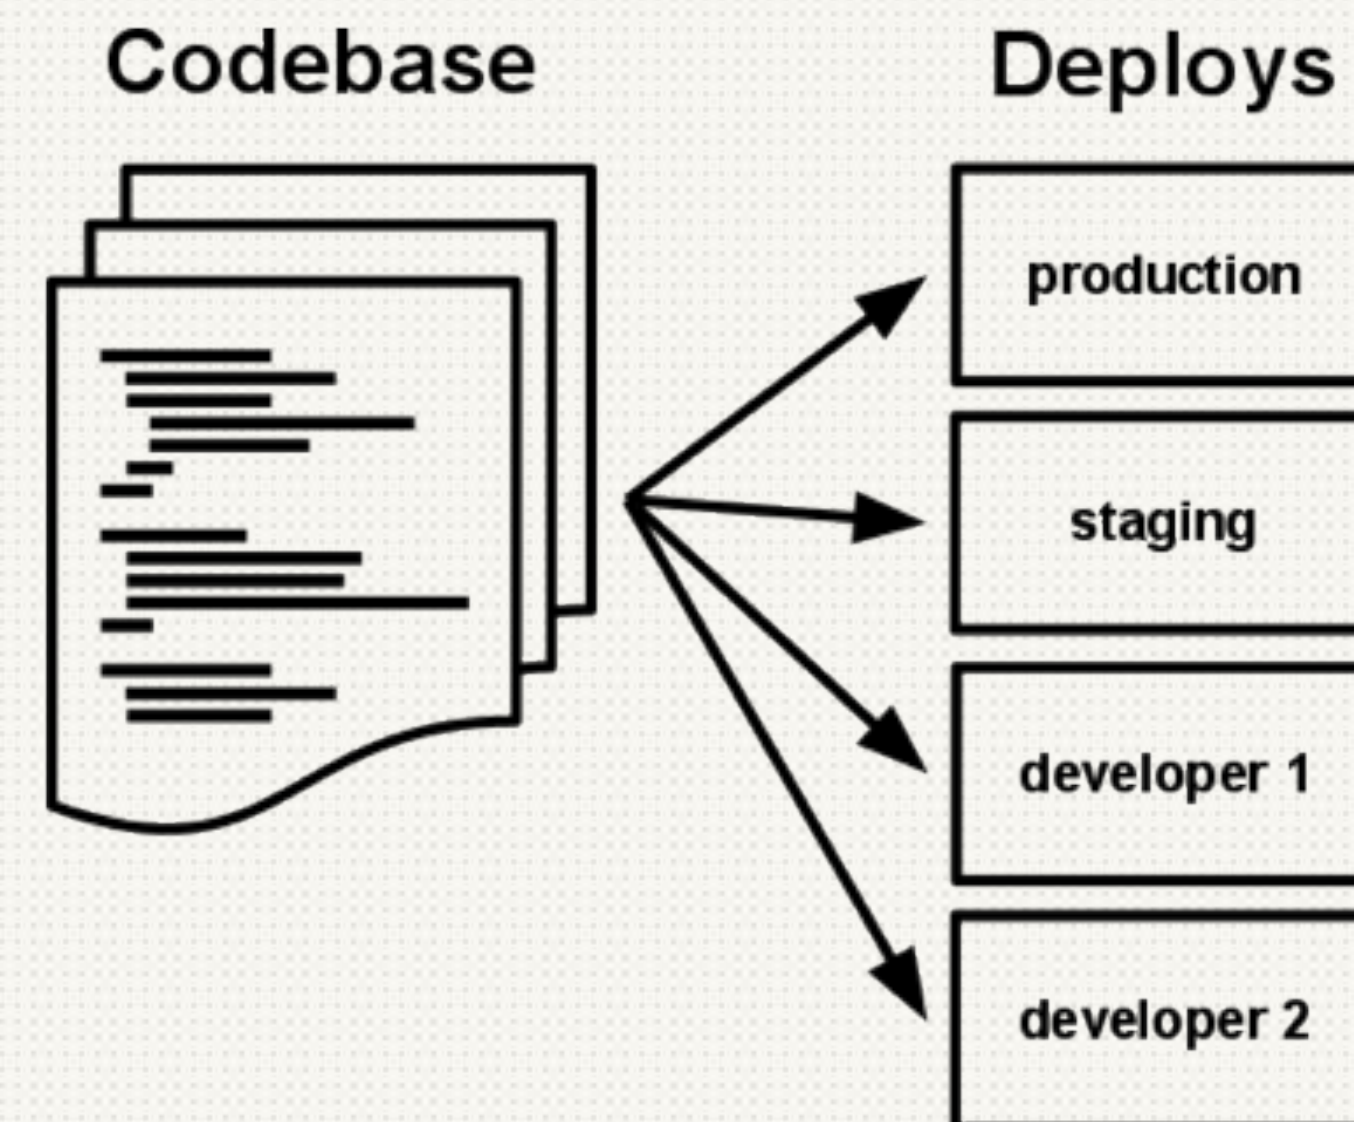
12FACTOR.NET

CODEBASE



I. CODEBASE

- ☑ 1 codebase, kept in a Version control system
- ☑ 1 codebase == 1 app. No cheating.
- ☑ Shared code is in libraries.
- ☑ Library sources are NOT included in app sources.
- ☑ One codebase can have many deploys (dev, test, staging, prod)





DEPENDENCIES

II. DEPENDENCY DECLARATION

- ▶ never rely on implicit existence of system-wide packages
 - ▶ e.g. curl does **not** exist everywhere your app will run!! (yes, docker, i'm talking to you)
- ▶ "declare all dependencies, completely and exactly, via a dependency declaration manifest"
 - ▶ completely: don't leave anything out. (this is not as easy as it sounds)
 - ▶ exactly: make sure to specify versions, or version upgrade policy
 - ▶ optimistic `>=` vs. pessimistic `~>` version constraints (semantic versioning)

II. DEPENDENCY ISOLATION


- ◆ Use some environment that allows you to ensure you are loading what you've declared.
 - ◆ And only that!
 - ◆ if your stack doesn't provide it, just start with as empty a box as you can find
- ◆ Examples:
 - ◆ Ruby: gemfile + bundle exec
 - ◆ Python: pip + virtualenv
 - ◆ Java: classpath

CONFIG



III. CONFIG

- ▶ IS: Everything that varies between deploys.
- ▶ (db url, credentials, hostnames, etc.)
- ▶ NOT constants in code. (Config varies substantially across deploys, code does not.)
- ▶ Avoid config files.
- ▶ All config in env vars.




BACKING SERVICES

IV. BACKING SERVICES

- Database. Obviously.
- Message queue (rabbitmq, SQS, redis, etc.)
- SMTP client (postfix)
- Cache (memcached, redis)
- Metrics gathering (new relic, loggly)
- Binary object stores (Amazon S3)

IV. BACKING SERVICES

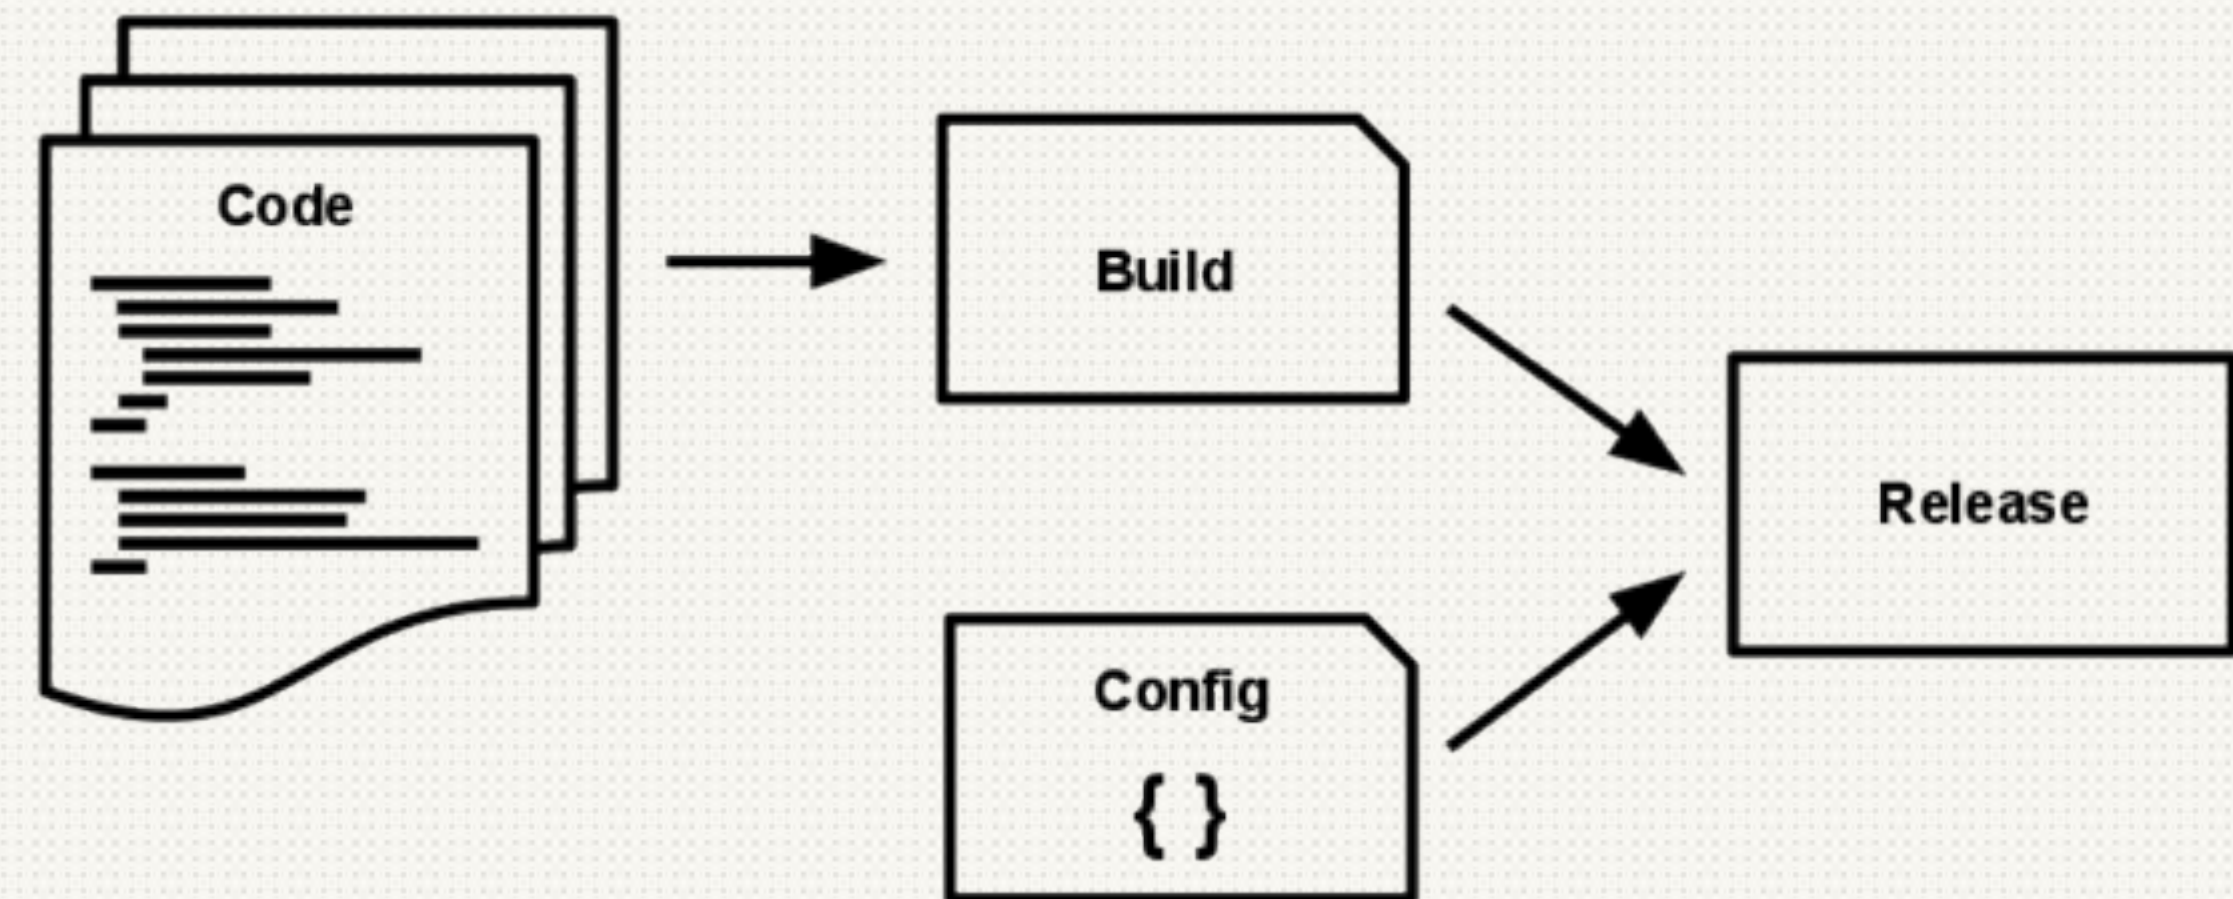
- ◆ Make no distinction between local and third-party services
 - ◆ `localhost` shouldn't make a difference to `some_fqdn`
- ◆ Should be easy to swap simply by changing resource handle (URL) and credentials.
 - ◆ e.g. locally hosted mysql / postgresql <==> AWS RDS
- ◆ Attachable resources: i am connecting to some db, when that is available.



BUILD – RELEASE – RUN

V. BUILD – RELEASE – RUN


- * Strictly separate the below 3 build and run stages:
- * build: create executable bundle, using some version, fetching all dependencies and compiling (where necessary).
- * release: add config to bundle, to have something that's ready to run.
- * run: (a.k.a. runtime) launch some process(es) that run the release.



V. BUILD – RELEASE – RUN

- * No changes to code at runtime.
- * Use deployment tools for release management:
 - * deploy automatically,
 - * but also ****rollback**** a release.
- * Releases are immutable.
 - * Yes, this means you are not allowed to manually make changes to production!
- * Releases should have some unique ID
 - * timestamp, number increment, hash, etc.

PROCESSES



VI. PROCESSES

- ☑ Execute the app as one or more stateless processes.
- ☑ Twelve-factor processes are stateless and use a shared-nothing architecture.
- ☑ Data is persisted in stateful backing services, typically a database.
 - ☑ Use memory or local filesystem ONLY as a brief, single-transaction cache.
 - ☑ Don't expect the filesystem or memory to be the same.

PORT BINDING

VII. PORT BINDING

- ▶ Export services via port binding
- ▶ Self-contained services that listen to a specific port.
- ▶ webserver library is bundled into the app. Do not rely on server existing in system.



CONCURRENCY

VIII. CONCURRENCY

- ▶ Separate Workload type \leftrightarrow process type.
 - ▶ (e.g. HTTP requests \leftrightarrow web process, long-running background tasks \leftrightarrow worker process.)
- ▶ Many workloads, but also many processes of the same workload!
- ▶ Rely on some process manager or scheduler to handle lifecycle of your processes (don't daemonize!)

DISPOSABILITY

IX. DISPOSABILITY

- Process must start (and STOP!!) fast!
- Minimize startup time (a few seconds!)
- Shutdown gracefully (on SIGTERM)
- Protect your state against sudden death.

DEV / PROD PARITY

X. DEV / PROD PARITY

	Traditional app	Twelve-factor app
Time between deploys	Weeks	Hours
Code authors vs code deployers	Different people	Same people
Dev vs production environments	Divergent	As similar as possible

X. DEV/PROD PARITY

- * Keep development, staging, and production as similar as possible
- * Use same backing services (same type of) between development and production
 - * e.g. using in-memory dbs is an anti-pattern
- * Design for continuous deployment

LOGS



XI. LOGS

- Use logs as event streams
 - a stream of aggregated, time-ordered events from all running processes and backing services.
- Not concerned with routing or storage of output stream.
 - Do not attempt to write to or manage logfiles (e.g. rotate). Instead, each running process writes its event stream, unbuffered, to stdout.
- In production, a log router / forwarder will take care of sending events from log stream to a centralized logging facility for long-term storage.
- where... find events (fast!), spot patterns (graphs over time), alerting!

ADMIN PROCESSES

XII. ADMIN PROCESSES

- ▶ one-off tasks:
 - ▶ in an identical environment as long-running processes,
 - ▶ running against a release,
 - ▶ using same codebase and
 - ▶ using same config.
- ▶ generally shipped together with production code, using same dependency isolation.
 - ▶ they become part of the release



That's all Folks!

THANK YOU!

Yorgos Saslis
@gsaslis
github.com/gsaslis