


Intro to Concurrency in Go

Ioannis Sermetziadis
Backend Engineer

June 26, 2020

Roadmap

Introduction to Go
Introduction to Concurrency
Concurrency + Go
Why Go?



Introduction to Go

Language history

- ▶ Design started in 2007 @ Google
- ▶ Robert Griesemer, Rob Pike, and Ken Thompson, Ian Lance Taylor and Russ Cox
- ▶ 1.0 version on 03/2013, 6-month release cycle
- ▶ Latest version 1.14 on 02/2020
- ▶ Open sourced since 2009
- ▶ Inspired from paper Communicating Sequential Processes (Hoare, 1978)
- ▶ Concurrency features of other languages, like Limbo (Bell Labs)

Core values

- ▶ Programming productivity for multi-core, networked, large systems
- ▶ Code readability - required on large and complex codebases
- ▶ Simplicity / Minimalism
- ▶ Go 1 is fixed to language specification
- ▶ Address criticism of other languages by combining best-practices
- ▶ Solving the painpoints of large systems, like build time, dependency management, readability

Characteristics

- ▶ Static typing
- ▶ Garbage collection
- ▶ Modular - versioned modules and dependency management
- ▶ Cross-compile from / to different platforms using \$GOOS and \$GOARCH
- ▶ Provides binaries for freeBSD, linux, macOS and windows

Hello World!

```
package main

import (
    "fmt"
)

func main() {
    fmt.Print("Hello DevStaff!")
}
```

Hello panic!

```
package main

import "time"

func main() {
    go func() {
        panic("failed!")
    }()

    time.Sleep(time.Second)
}
```


Hello HTTP!

- ▶ applies implicit concurrency
- ▶ a goroutine triggered per HTTP connection

```
func helloHandler(w http.ResponseWriter, req *http.Request) {  
    fmt.Fprintf(w, "hello DevStaffer!\n")  
}
```

```
func main() {  
    log.Println("listening...")  
    http.HandleFunc("/", helloHandler)  
    err := http.ListenAndServe("localhost:12345", nil)  
    if err != nil {  
        log.Fatal("ListenAndServe: ", err)  
    }  
}
```

HTTP client

- ▶ HTTP client is thread-safe - no extra effort to support calls by multiple goroutines

```
resp, err := http.Get("http://localhost:12345")
if err != nil {
    log.Fatalln(err)
}

defer resp.Body.Close()

body, err := ioutil.ReadAll(resp.Body)
if err != nil {
    log.Fatalln(err)
}

log.Println(fmt.Sprintf("received: %s", string(body)))
```

Introduction to Concurrency

- ▶ Concurrency abstracts independently executing workloads and improves responsiveness
- ▶ "out of order" / non-deterministic execution that still produces consistently the same result
- ▶ Can make programs faster or slower
- ▶ Sequential thinking and sequential design are simpler, but don't model real-world problems
- ▶ Async callbacks are not easy to read and understand

Concurrency vs Parallelism

- ▶ Concurrency = when two or more tasks start, run and complete in overlapping time periods
- ▶ Tasks need to be interruptable
- ▶ Parallelism = when two or more tasks run at the same time
- ▶ Tasks need to be independent

CPU or IO bound

- ▶ CPU-bound workloads require parallelism because CPU cores are always busy
- ▶ Not all CPU-bound workloads (algorithms) are suitable for concurrency - its expensive to break-up work or combine results
- ▶ IO-bound workloads don't need parallelism - more threads than cores can improve workload execution
- ▶ IO-bound workloads are suitable for concurrency

OS Scheduler

- ▶ Software component responsible for allocating CPU time to processes or threads
- ▶ A Process can spawn many Threads allocated to 1 or more CPU cores
- ▶ Context-switch happens when a Thread is deallocated from a CPU core and another Thread is allocated
- ▶ Thread states - Waiting, Runnable, Executing
- ▶ Preemptive scheduling - time sliced scheduling - preempt/interrupt and resume

Concurrency in Go

goroutines (**go**)
channels (**chan**)
select (**select**)

Goroutine

- ▶ Goroutine also known as light-weight process or green thread
- ▶ Derived from coroutines and inspired by the Communicating Sequential Processes paper

Goroutines example

```
const jobs = 100000

func main() {
    defer logSpan(time.Now())

    // WaitGroup synchronises on a collection of goroutines signal
    wg := sync.WaitGroup{}
    wg.Add(jobs)

    for i := 0; i < jobs; i++ {
        // spawn a go routine for each job
        go func() {
            // simulate some work
            time.Sleep(1 * time.Second)
            // signal finish
            wg.Done()
        }()
    }

    log.Println("waiting jobs to complete...")
    wg.Wait()
}
```

Channel

- ▶ Enable goroutine communication in blocking or non-blocking way
- ▶ Typed primitives that support send and receive operations and allow goroutines to **synchronise** and **communicate**
- ▶ Operations are blocking by default (unbuffered)
- ▶ Non-blocking send operation is possible by defining channel buffer capacity
- ▶ Goroutines can synchronize on channels without explicit locking, by blocking on send and receive
- ▶ Wired into the language API (e.g. time package)

Channel operations

- ▶ Receive operation using `for`, `select` or simple receive expression
- ▶ Receiver can query the state of the channel, using an extra return value '`v, ok := <- ch`'
- ▶ A closed channel will return zero values (of channel type) on receive operation
- ▶ Send operation to a closed channel will trigger panic - sender should close the channel when it stops sending
- ▶ Channel directions increase type safety by statically checking the operation
- ▶ Write will block when buffer is full and read when buffer is empty

Producer-Consumer example

```
func produce(msgs chan<- msg) {  
    for {  
        w := rand.Intn(100)  
        msgs <- msg{  
            value: w,  
        }  
        time.Sleep(time.Duration(w) * time.Millisecond)  
    }  
}  
  
func consume(msgs <-chan msg) {  
    for msg := range msgs {  
        fmt.Println(msg.value)  
    }  
}  
  
func main() {  
    var msgs = make(chan msg)  
    go produce(msgs)  
    go consume(msgs)  
    time.Sleep(2 * time.Second)  
}
```

Select

- ▶ Control statement similar to switch, but decision is made based on the ability to communicate
- ▶ Goroutines can operate (send or receive) on multiple channels
- ▶ If possible to operate on multiple channels, chooses pseudo-randomly
- ▶ **default** clause provides safety for immediate outcome

Select example

```
func main() {  
    input := make(chan int)  
    output := make(chan int)  
  
    select {  
    case input <- 1:  
    case _ = <-output:  
        // case <-time.After(time.Second):  
        //      log.Println("timed out")  
    default:  
        log.Println("no one was ready")  
    }  
}
```

Spadtest in Go

```
func main() {  
    for i := 0; i < 10; i++ {  
        var winner string  
        select {  
        case winner = <-ping("Google"):  
        case winner = <-ping("Facebook"):  
        case winner = <-ping("Twitter"):  
        case winner = <-ping("LinkedIn"):  
        case <-time.After(300 * time.Millisecond):  
            log.Println("none responded in time")  
            continue  
        }  
        log.Printf("got response from %s\n", winner)  
    }  
}
```

Why Go?

Go Scheduler

- ▶ Non-preemptive/co-operative scheduler that depends on user-space events used for scheduling decisions
- ▶ Part of the Go runtime built into each Go application
- ▶ Prevents thread context-switches by allocating work
- ▶ Uses Global Run Queue (GRQ) and Local Run Queue (LRQ) to keep track of the goroutines state and dynamically assign to Threads
- ▶ Goroutines are context-switched on and off a Thread, similar to Threads context-switched on a CPU core
- ▶ When a Goroutine is blocked (system call, IO, synchronisation), it gets de-allocated and moved to the network poller

OS vs Go scheduler

- ▶ Work Stealing - Goroutines are balanced across all available threads to avoid OS context-switches
- ▶ When a thread LRQ is empty, scheduler steals work from other thread LRQ
- ▶ In C, OS thread context-switching costs ~ 1000 nanoseconds or $\sim 17k$ instructions - cache misses also possible
- ▶ In Go, Goroutine context-switching prevents thread content-switching - cost is ~ 150 nanoseconds or $\sim 2.4k$ instructions
- ▶ The Go scheduler turns the IO-bound workload into CPU-bound work that leverages CPU capacity
- ▶ Preemptive vs Non-preemptive scheduling - timesliced execution vs cooperation

Data race detection

- ▶ Multiple goroutines access the same variable and at least one is a write
- ▶ The result of a program will vary randomly based on the order to access

Data race examples:

- ▶ Read of loop counter variable from child goroutines
- ▶ Accidentally shared variable
- ▶ Unprotected global variable

```
func main() {  
    for i := 0; i < 1000; i++ {  
        state := 10  
        wg := sync.WaitGroup{}  
        wg.Add(2)  
        go multiply(&state, &wg, 10)  
        go deduct(&state, &wg, 2)  
  
        wg.Wait()  
        fmt.Println(state)  
    }  
}
```

Deadlock errors

- ▶ A group of goroutines are waiting for each other and none of them is able to proceed
- ▶ Possible in message-passing or shared-memory concurrency models
- ▶ Go detects only when the whole program freezes not for subset of goroutines
- ▶ Channel deadlocks are easier to detect than Lock deadlocks

```
func main() {  
    ch := make(chan int)  
    ch <- 1  
    fmt.Println(<-ch)  
}
```

Links

Tour of Go

tour.golang.org/welcome/1

Playground

play.golang.org/

Go Concurrency Patterns

talks.golang.org/2012/concurrency.slide

Advanced Go Concurrency Patterns

blog.golang.org/io2013-talk-concurrency

Scheduling In Go

www.ardanlabs.com/blog/2018/08/scheduling-in-go-part1.html